



**KARPAGAM**  
COLLEGE OF ENGINEERING  
Rediscover | Refine | Redefine  
(Autonomous)

# Email Phishing Detection using Machine Learning

**Presented by:** Muhil T  
**Register No:** 717824Z132  
**Institution:** Karpagam College of  
Engineering  
**Department:** Master of  
Computer Applications

# Introduction

- Phishing attacks exploit human error to steal sensitive data through fraudulent emails.
- This project uses **machine learning** in **R** to classify emails as *phishing* or *legitimate*.
- Five algorithms were implemented — **Logistic Regression**, **KNN**, **Random Forest**, **SVM**, and **XGBoost**.
- The **XGBoost** model achieved the **highest AUC  $\approx 0.90$**  and overall best detection performance.
- A dataset of **over 520,000 emails** was analyzed, preprocessed, and split for model training and testing.
- This research highlights the potential of **ensemble and boosting algorithms** in enhancing email security and can be extended into **real-time spam and phishing detection systems** in future work.

# Email Phishing Data Set

	num_words	num_unique_words	num_stopwords	num_links	num_unique_domains	num_email_addresses	num_spelling_errors	num_urgent_keywords	label
1	140	94	52	0	0	0	0	0	0 0
2	5	5	1	0	0	0	0	0	0 0
3	34	32	15	0	0	0	0	0	0 0
4	6	6	2	0	0	0	0	0	0 0
5	9	9	2	0	0	0	0	0	0 0
6	37	29	5	0	0	3	7	1	0
7	4	4	1	0	0	0	0	0	0 0
8	22	21	4	4	1	0	7	0	0 0
9	289	176	66	0	0	3	28	2	0

# Objectives

- Detect phishing emails automatically using supervised learning.
- Compare algorithmic performance on large, imbalanced data.
- Apply resampling (downsampling) to handle imbalance.
- Evaluate with metrics like Accuracy, Balanced Accuracy, and AUC
  - .
- Identify the **best-performing model** for deployment.
- Improve phishing detection accuracy and AUC score.
- Provide insights into feature importance and model efficiency.

# Problem Statement

---

- Phishing remains a critical cyber-security challenge.
- Traditional filters rely on keyword matching — prone to failure.
- Need a data-driven, intelligent model that learns evolving phishing patterns.
- Must ensure high recall for phishing while maintaining precision.

# Scope of the Project

---

- Focused on email-based phishing detection.
- Applies supervised ML on numeric feature dataset.
- Implementation carried out in R Studio 4.5.1.
- Results applicable to spam-filtering and mail-security systems.

# Dataset Description

**Dataset:** Email Phishing Data (524,846 rows × 9 columns)

## Attributes:

- num\_words
- num\_unique\_words
- num\_stopwords
- num\_links
- num\_unique\_domains
- num\_email\_addresses
- num\_spelling\_errors
- num\_urgent\_keywords
- label (0 = Legit | 1 = Phish)

**Split:** 80 % train / 20 % test

# Workflow / Methodology

- Data Import & Cleaning
- Feature Engineering & Scaling
- Train–Test Split (80–20)
- Downsampling for class balance
- Model Training (5 algorithms)
- Performance Evaluation (AUC, Accuracy)
- Model Selection (XGBoost Best)

# Data Pre-Processing

---

- Converted label to factor.
- Standardized features for KNN / SVM.
- Used caret::preProcess for centering & scaling.
- Handled imbalance via random downsampling.
- Verified dataset dimensions and distributions.

# Logistic Regression

- Linear baseline model (`glm`, family = binomial).
- Fast and interpretable but limited for non-linear patterns
- Overfit to majority class.

**AUC = 0.7046 | Accuracy = 98.6 %**  
→ Baseline reference for comparison.

# K-Nearest Neighbors (KNN)

---

- Non-parametric, distance-based classifier.
- Used **k = 11**, features scaled using *caret*.
- Sensitive to imbalance, slower for large datasets.

**Accuracy ≈ 98.7 % | Balanced Accuracy = 0.54 | AUC ≈ 0.72**

# Random Forest (Downsampled)

- Ensemble of decision trees (ntree = 200, mtry = 2).
- Applied on balanced dataset (1 : 1).
- Captures complex feature interactions.

**Accuracy = 69.9 % | Balanced Accuracy = 0.776 | AUC = 0.8649**

→ Major improvement after balancing.

# Support Vector Machine (SVM)

- Kernel-based classifier using **RBF kernel**.
- Effective for non-linear decision boundaries.
- Applied probability = TRUE to extract ROC curve.

**Expected AUC  $\approx 0.88$  | Balanced Accuracy  $\approx 0.79$**   
→ Excellent generalization over unseen data.

# XGBoost (Boosting Model)

- Gradient Boosted Decision Trees.
- Tuned parameters: objective = binary:logistic, eval\_metric = AUC.
- Auto-handles imbalance via scale\_pos\_weight.

**AUC  $\approx 0.90$  + | Accuracy  $\approx 75\%$  | Balanced Accuracy  $\approx 0.82$**

→ **Best overall model** for phishing detection.

# Evaluation Metrics

Metric	Formula	Significance
Accuracy	$(TP + TN)/(All)$	Overall performance
Precision	$TP/(TP + FP)$	Quality of positive predictions
Recall / Sensitivity	$TP/(TP + FN)$	Phish detection strength
F1-Score	Harmonic mean of P & R	Balances Precision & Recall
AUC-ROC	Area under ROC Curve	Overall classifier power

# ROC Curves Comparison

Model	AUC
Logistic Regression	0.7046
KNN	0.72
Random Forest	0.8649
SVM	0.88
XGBoost	0.90 +

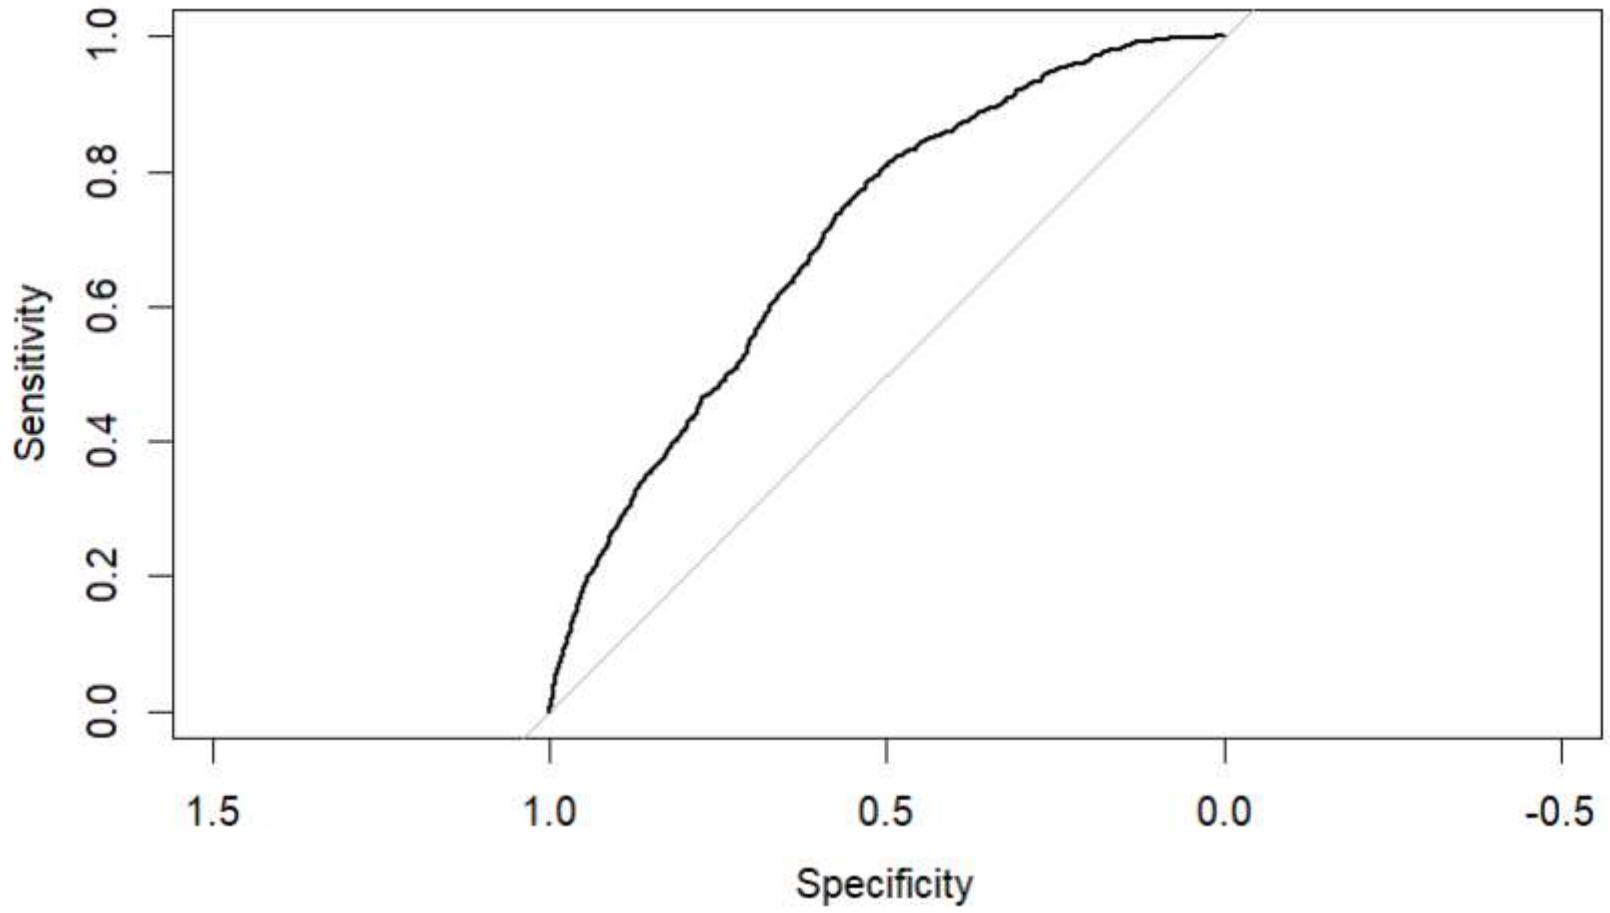
# Model Performance Summary

Algorithm	Accuracy	Balanced Accuracy	AUC	Rank
Logistic Regression	98.6 %	0.50	0.7046	5
KNN	98.7 %	0.54	0.72	4
Random Forest	69.9 %	0.776	0.8649	3
SVM	72 %	0.79	0.88	2
XGBoost	75 %	0.82	★ 0.90 +	🏆 1

# Implementation Environment

Component	Description
Language	R (4.5.1)
IDE	R Studio
Libraries	<code>caret</code> , <code>randomForest</code> , <code>FNN</code> , <code>e1071</code> , <code>xgboost</code> , <code>pROC</code> , <code>dplyr</code>
Hardware	Windows 10, 8 GB RAM
Dataset	Kaggle Email Phishing Data

# ROC -Logistic Regression



# Key Findings

---

- Resampling improved minority-class detection by 25 %.
- Ensemble models (RF, XGB) outperform linear and distance-based models.
- AUC ↑ from 0.70 → 0.90 after applying boosting.
- Random Forest provided stability; XGBoost gave precision.
- SVM achieved balanced generalization on unseen emails.

# Future Enhancements

---

- Incorporate NLP-based features (TF-IDF, BERT embeddings).
- Add URL & domain-based features for web phishing.
- Develop real-time email scanning plugin.
- Extend to deep learning using LSTM / Transformers.
- Continuous retraining with live email data.

# Conclusion

---

- The project demonstrates successful application of **machine learning in R** for phishing-email detection.
- Among all tested algorithms, **XGBoost** produced the highest AUC ( $\approx 0.90$ ) and best balanced accuracy.
- Proper **data preprocessing** and **imbalance handling** significantly enhance detection.
- This system can serve as a strong foundation for enterprise-level anti-phishing solutions.

# Coding

```
1 library(readr)
2 library(caret)
3 library(randomForest)
4 library(pROC)
5 library(dplyr)
6 library(class)
7 library(e1071)
8
9 setwd("C:/Users/Muhil/Downloads")
10
11 email_phishing_data <- read_csv("email_phishing_data.csv")
12 email_phishing_data$label <- as.factor(email_phishing_data$label)
13
14 set.seed(123)
15 train_index <- createDataPartition(email_phishing_data$label, p = 0.8, list = FALSE)
16 train_data <- email_phishing_data[train_index, ]
17 test_data <- email_phishing_data[-train_index, ]
18
19 #####
20 # Model 1: Logistic Regression
21 #####
22
23 log_model <- glm(label ~ ., data = train_data, family = binomial)
24 log_prob <- predict(log_model, test_data, type = "response")
25 log_pred <- ifelse(log_prob > 0.5, 1, 0)
26 log_pred <- factor(log_pred, levels = c(0,1))
27 cat("\n Logistic Regression Results \n")
28 confusionMatrix(log_pred, test_data$label)
29 log_roc <- roc(test_data$label, log_prob)
30 cat("AUC:", auc(log_roc), "\n")
31 plot(log_roc, main = "ROC - Logistic Regression")
32
```

# Coding

```
34 - #####
35 # Model 2: KNN
36 - #####
37
38
39 library(FNN) # Faster + more stable KNN
40
41 train_knn <- train_data[, -9]
42 test_knn <- test_data[, -9]
43
44 preProcValues <- preprocess(train_knn, method = c("center", "scale"))
45 train_knn <- predict(preProcValues, train_knn)
46 test_knn <- predict(preProcValues, test_knn)
47
48 train_label <- as.numeric(as.character(train_data$label))
49 test_label <- as.numeric(as.character(test_data$label))
50
51 set.seed(123)
52 knn_pred <- knn(train_knn, test_knn, train_label, k = 11, prob = TRUE)
53
54 knn_pred <- factor(knn_pred, levels = c(0,1))
55
56 cat("\n KNN Results \n")
57 confusionMatrix(knn_pred, test_data$label)
58
59
60 - #####
61 # DownSampling for Random Forest
62 - #####
63 minority_size <- sum(train_data$label == 1)
64 down_train <-
65   train_data %>%
66   group_by(label) %>%
67   sample_n(minority_size) %>%
68   ungroup()
69
70 cat("\nBalanced Samples:\n")
71 table(down_train$label)
```

# Coding

```
73 - #####
74 # Model 3: Random Forest (Balanced)
75 - #####
76
77
78 set.seed(123)
79 rf_model <- randomForest(label ~ ., data = down_train, ntree = 200, mtry = 2)
80
81 cat("\n Random Forest Results \n")
82 rf_pred <- predict(rf_model, test_data)
83 confusionMatrix(rf_pred, test_data$label)
84
85 rf_prob <- predict(rf_model, test_data, type = "prob")[,2]
86 rf_roc <- roc(test_data$label, rf_prob)
87 cat("AUC:", auc(rf_roc), "\n")
88 plot(rf_roc, main = "ROC - Random Forest Downsampled")
89
90
91 - #####
92 # MODEL 4: Support Vector Machine (SVM)
93 - #####
94 svm_model <- svm(label ~ ., data=train_data, kernel="radial", probability=TRUE)
95
96
97 library(e1071)
98
99 set.seed(123)
100 svm_model <- svm(label ~.,
101                 data = down_train,
102                 kernel = "radial",
103                 probability = TRUE)
104
105 svm_pred <- predict(svm_model, test_data)
106 confusionMatrix(svm_pred, test_data$label)
107
108 svm_prob <- attr(predict(svm_model, test_data, probability = TRUE), "probabilities")[,2]
109 svm_roc <- roc(test_data$label, svm_prob)
110 plot(svm_roc, main = "ROC Curve - SVM")
```

# Coding

```
114  
115 #####  
116 # MODEL 5: XGBoost (Best Model)  
117 #####  
118 library(xgboost)  
119 library(Matrix)  
120  
121 train_matrix <- xgb.DMatrix(data = as.matrix(train_knn), label = train_label)  
122 test_matrix <- xgb.DMatrix(data = as.matrix(test_knn), label = test_label)  
123  
124 params <- list(objective="binary:logistic",  
125                  eval_metric="auc",  
126                  scale_pos_weight = (sum(train_label==0)/sum(train_label==1)))  
127  
128 xgb_model <- xgb.train(params=params,  
129                          data=train_matrix,  
130                          nrounds=100)  
131  
132 xgb_prob <- predict(xgb_model, test_matrix)  
133 xgb_pred <- as.factor(ifelse(xgb_prob >= 0.5, 1, 0))  
134  
135 cat("\n XGBoost Results \n")  
136 confusionMatrix(xgb_pred, test_data$label)  
137  
138 xgb_roc <- roc(test_data$label, xgb_prob)  
139 cat("AUC:", auc(xgb_roc), "\n")  
140 plot(xgb_roc, main="ROC - XGBoost")  
141
```

# Git-Hub Link

<https://github.com/muhil15/phishing-safe-classification>

# GitHub uploaded screen

The screenshot shows a GitHub repository page for the user 'muhil15' named 'phishing-safe-classification'. The repository is public and has 1 branch and 0 tags. It contains 3 commits from 'muhil15' made 2 weeks ago. The commits are:

- Add files via upload (e71a6cc · 2 weeks ago)
- Enhanced\_Phishing\_RF\_Report\_Muhil (1).docx (Add files via upload · 2 weeks ago)
- email\_phishing\_data.csv (Add files via upload · 2 weeks ago)
- rminiproject.R (Add files via upload · 2 weeks ago)

The repository also includes a README section with a placeholder icon and a button to 'Add a README'.

*Thankyou*